


Whale AI — grounded, tiered, and inside the perimeter

50+ use cases across every family, on your choice of model, including fully offline



AI in every family. Inside the perimeter.

September 2026 · Version 1.0 · www.bluewhalestack.com

How Whale AI grounds every answer in live inventory, billing, tickets and findings, routes work across the Spark, Tide and Abyss tiers, and runs offline for sovereign estates.

Whale AI is an intelligence layer woven into every module rather than a chatbot beside them: 50+ production use cases across cost, security, ITSM, operations, migration and governance, each grounded in your live inventory, billing, tickets and findings. It runs in three tiers — Spark, Tide and Abyss — with your choice of model, including fully offline inside the perimeter for sovereign estates.



50+ Use cases, each with declared data dependencies	Spark · Tide · Abyss Three tiers routed by task complexity	Offline In-region and air-gapped models for sovereign deployments
---	--	---

How it is built

Grounded, not generative guesswork Every use case declares the providers it reads — inventory, cost, tickets, findings — and cites them in the answer; numbers come from the platform, not the model.	Tiered intelligence Spark answers interactive questions in under three seconds; Tide reasons over migration plans and IaC; Abyss runs multi-step agents for the hardest cross-module work.
In the flow of work An assistant drawer in every module, streamed answers, exports to PDF and Excel, and prompt caching so repeat questions are instant.	Inside the perimeter Bring your model or run ours in-region; sovereign deployments run Whale AI fully offline with zero external data exposure.

From question to grounded answer

- 1

Ask, in context
From any screen, ask about what you are looking at — a workload, a bill, an incident.
- 2

Ground
The use case gathers the live facts it declares, scoped to your permissions.
- 3

Answer with citations
The right tier reasons over the facts and streams an answer with the evidence behind it.

Three tiers

Spark — Standard & up Entry, fast — summaries, inventory Q&A and sizing recommendations in natural language.	Tide — Enterprise & up Deep reasoning for migration planning, IaC drafting and multi-cloud analysis.
Abyss — Enterprise · Government (in-region, offline) Premium multi-step agents for the most complex, cross-module workflows — including fully offline inside the perimeter.	

Use cases by category

Category	Use cases
FinOps & cost	12
Security	10
ITSM & incidents	8
Cloud operations	8
Infrastructure	6
Observability	5
Executive intelligence	4
Governance	4
Migration	3
Performance	3

Counts as shipped in the September 2026 release; every use case declares the data providers it reads and cites them in the answer.

Grounding providers

- resource_inventory — the live resource model and workloads
- cost_summary and cost_recommendations — billing joined to inventory, rightsizing and commitments
- scanner_findings — security and compliance findings
- open_incidents — ITSM incidents and changes
- tenant_metadata — tenant scope and permissions

Where your data goes

On SaaS, requests are processed in your chosen region on the hosted model you select. On BYOC they run in your accounts. On sovereign and air-gapped deployments Whale AI runs fully offline on in-region models with zero external calls. Every call is scoped to the authenticated tenant, with isolation enforced at the database layer.

Questions security and data teams ask

01 Does our data leave our environment?

On SaaS, requests are processed in your chosen region; on BYOC they run in your accounts; on sovereign deployments Whale AI runs fully offline on in-region models with no external calls.

02 Which models are used?

Your choice — hosted models including Anthropic Claude, or models you bring and run inside the perimeter. Use cases are model-agnostic; grounding and citations work the same way.

03 Can it take actions?

Use cases draft plans, reports and configurations; actions still go through the platform's approvals and audit trail, so Whale AI never bypasses governance.

© 2026 BlueWhale Stack. Product, edition and maturity statements are as of September 2026; certifications are described exactly as their issuing bodies allow. Customer engagements are anonymized under confidentiality. Contact:

sales@bluewhalestack.com.



BLUEWHALE STACK

Every cloud, one control plane — unified inventory, AI-native provisioning, observability, migration, and governance across public cloud, on-prem, and hybrid.

CLOUD INSIGHTS, MONTHLY

Subscribe

Product

Platform overview

Modules

Whale AI

WhaleForge IaC

Editions & pricing

BlueWhale Stack Fabric

Solutions

By outcome

Government

BFSI

Healthcare

Telco & MSP

Datacenter & Colocation

Customer stories

Case studies

Developers

Documentation

Quick start

Cloud integration

Identity & access

API reference

Partner Portal

Company

About us

Leadership

Careers

Partners

Newsroom

Resources

Trust Center

Contact

ISO 27001:2022

ISO 27017 (cloud security)

ISO 27018 (PII in cloud)

ISO 27701 (privacy ISMS)

ISO 22301 (business continuity)

CSA STAR Level 1

SOC 2 Type II (readiness)

GDPR

India DPDP Act 2023

© 2026 BlueWhale Stack. BlueWhale Stack is a registered trademark. All rights reserved.

[Privacy](#) [Terms](#) [LinkedIn](#)